

R Notebook

##Introduction

a) Specifying the Question

The main objective of the study is to use the data provided to aid Carrefour's marketing team in formulating strategies to boost sales.

b) Defining the Metrics for Success

- Determining and visualising the descriptive statistics of the variables in the datasets provided.
- Carrying out principal component analysis.
- Carrying out feature selection.
- Performing association analysis.
- Identifying possibly fraudulent transactions.

c) Understanding the context

Sales and Marketing teams aim to maximise a business' profit. Data-driven insights allow for the planning of more targeted and effective campaigns.

d) Recording the Experimental Design

- Determine the main objectives.
- Load and preview the datasets.
- Understand the data.
 - Prepare the datasets - Identify outliers, anomalies, duplicates, missing values, and determine how deal with them, drop unnecessary columns etc.
 - Analyse the data using univariate, bivariate, and multivariate analysis techniques.
 - Carry out dimensionality reduction, feature selection, associative analysis and anomaly detection on the respective datasets
- Conclusion and recommendations

e) Data Relevance

The datasets provided were relevant to the research question as they had relevant details on the sales at Carrefour.

```
##Loading the dataset
```

```
#Loading some required libraries
```

```
library(readr)
```

```
library(data.table)
```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(psych)
```

```
##
```

```
## Attaching package: 'psych'
```

```
## The following objects are masked from 'package:ggplot2':
```

```
##
```

```
##      %+%, alpha
```

```
library(Rtsne)
```

```
library(tidyverse)
```

```
## — Attaching packages
```

```
##
```

```
## tidyverse 1.3.2 —
```

```
## ✓ tibble 3.1.7      ✓ dplyr 1.0.9
```

```
## ✓ tidyr 1.2.0       ✓ stringr 1.4.0
```

```
## ✓ purrr 0.3.4       ✓ forcats 0.5.1
```

```
## — Conflicts ————— tidyverse_conflict
```

```
s() —
```

```
## ✗ psych::%+%( ) masks ggplot2::%+%( )
```

```
## ✗ psych::alpha( ) masks ggplot2::alpha( )
```

```
## ✗ dplyr::between( ) masks data.table::between( )
```

```
## ✗ dplyr::filter( ) masks stats::filter( )
```

```
## ✗ dplyr::first( ) masks data.table::first( )
```

```
## ✗ dplyr::lag( ) masks stats::lag( )
```

```
## ✗ dplyr::last( ) masks data.table::last( )
```

```
## ✗ purrr::lift( ) masks caret::lift( )
```

```
## ✗ purrr::transpose( ) masks data.table::transpose( )
```

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(ggbiplot)
```

```
## Loading required package: plyr
```

```
## -----
```

```

----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, th
en dplyr:
## library(plyr); library(dplyr)
## -----
----
##
## Attaching package: 'plyr'
##
## The following objects are masked from 'package:dplyr':
##
##     arrange, count, desc, failwith, id, mutate, rename, summarise,
##     summarize
##
## The following object is masked from 'package:purrr':
##
##     compact
##
## Loading required package: scales
##
## Attaching package: 'scales'
##
## The following object is masked from 'package:purrr':
##
##     discard
##
## The following objects are masked from 'package:psych':
##
##     alpha, rescale
##
## The following object is masked from 'package:readr':
##
##     col_factor
##
## Loading required package: grid

df <- fread("http://bit.ly/CarreFourDataset")
df <- data.frame(df)

```

##Checking the Data

Determining the no. of records in the dataset:

```

dim(df)

## [1] 1000    16

#the dataset has 1000 rows and 16 columns

```

Previewing the top of the dataset:

head(df)

##	Invoice.ID	Branch	Customer.type	Gender	Product.line	Unit.pric e
## 1	750-67-8428	A	Member	Female	Health and beauty	74.69
## 2	226-31-3081	C	Normal	Female	Electronic accessories	15.28
## 3	631-41-3108	A	Normal	Male	Home and lifestyle	46.33
## 4	123-19-1176	A	Member	Male	Health and beauty	58.22
## 5	373-73-7910	A	Normal	Male	Sports and travel	86.31
## 6	699-14-3026	C	Normal	Male	Electronic accessories	85.39

##	Quantity	Tax	Date	Time	Payment	cogs	gross.margin.percent age
## 1	7	26.1415	1/5/2019	13:08	Ewallet	522.83	4.761905
## 2	5	3.8200	3/8/2019	10:29	Cash	76.40	4.761905
## 3	7	16.2155	3/3/2019	13:23	Credit card	324.31	4.761905
## 4	8	23.2880	1/27/2019	20:33	Ewallet	465.76	4.761905
## 5	7	30.2085	2/8/2019	10:37	Ewallet	604.17	4.761905
## 6	7	29.8865	3/25/2019	18:30	Ewallet	597.73	4.761905

##	gross.income	Rating	Total
## 1	26.1415	9.1	548.9715
## 2	3.8200	9.6	80.2200
## 3	16.2155	7.4	340.5255
## 4	23.2880	8.4	489.0480
## 5	30.2085	5.3	634.3785
## 6	29.8865	4.1	627.6165

Previewing the bottom of the dataset:

tail(df)

##	Invoice.ID	Branch	Customer.type	Gender	Product.line	Unit.p rice
## 995	652-49-6720	C	Member	Female	Electronic accessories	60.95
## 996	233-67-5758	C	Normal	Male	Health and beauty	40.35
## 997	303-96-2227	B	Normal	Female	Home and lifestyle	97.38
## 998	727-02-1313	A	Member	Male	Food and beverages	3

```

1.84
## 999 347-56-2442 A Normal Male Home and lifestyle 6
5.82
## 1000 849-09-3807 A Member Female Fashion accessories 8
8.34
## Quantity Tax Date Time Payment cogs gross.margin.percenta
ge
## 995 1 3.0475 2/18/2019 11:40 Ewallet 60.95 4.7619
05
## 996 1 2.0175 1/29/2019 13:46 Ewallet 40.35 4.7619
05
## 997 10 48.6900 3/2/2019 17:16 Ewallet 973.80 4.7619
05
## 998 1 1.5920 2/9/2019 13:22 Cash 31.84 4.7619
05
## 999 1 3.2910 2/22/2019 15:33 Cash 65.82 4.7619
05
## 1000 7 30.9190 2/18/2019 13:28 Cash 618.38 4.7619
05
## gross.income Rating Total
## 995 3.0475 5.9 63.9975
## 996 2.0175 6.2 42.3675
## 997 48.6900 4.4 1022.4900
## 998 1.5920 7.7 33.4320
## 999 3.2910 4.1 69.1110
## 1000 30.9190 6.6 649.2990

```

Checking datatype of each column:

```

str(df)

## 'data.frame': 1000 obs. of 16 variables:
## $ Invoice.ID : chr "750-67-8428" "226-31-3081" "631-41-3108"
"123-19-1176" ...
## $ Branch : chr "A" "C" "A" "A" ...
## $ Customer.type : chr "Member" "Normal" "Normal" "Member" ...
## $ Gender : chr "Female" "Female" "Male" "Male" ...
## $ Product.line : chr "Health and beauty" "Electronic accessori
es" "Home and lifestyle" "Health and beauty" ...
## $ Unit.price : num 74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity : int 7 5 7 8 7 7 6 10 2 3 ...
## $ Tax : num 26.14 3.82 16.22 23.29 30.21 ...
## $ Date : chr "1/5/2019" "3/8/2019" "3/3/2019" "1/27/20
19" ...
## $ Time : chr "13:08" "10:29" "13:23" "20:33" ...
## $ Payment : chr "Ewallet" "Cash" "Credit card" "Ewallet"
...
## $ cogs : num 522.8 76.4 324.3 465.8 604.2 ...
## $ gross.margin.percentage: num 4.76 4.76 4.76 4.76 4.76 ...
## $ gross.income : num 26.14 3.82 16.22 23.29 30.21 ...

```

```
## $ Rating          : num  9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ Total           : num  549 80.2 340.5 489 634.4 ...
```

##Tidying the Dataset

#checking column names

```
colnames(df)
```

```
## [1] "Invoice.ID"          "Branch"
## [3] "Customer.type"       "Gender"
## [5] "Product.line"        "Unit.price"
## [7] "Quantity"            "Tax"
## [9] "Date"                "Time"
## [11] "Payment"             "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "Rating"              "Total"
```

#converting column names to lowercase

```
colnames(df) = tolower(colnames(df))
```

```
colnames(df)
```

```
## [1] "invoice.id"          "branch"
## [3] "customer.type"       "gender"
## [5] "product.line"        "unit.price"
## [7] "quantity"            "tax"
## [9] "date"                "time"
## [11] "payment"             "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "rating"              "total"
```

#checking for missing values

```
data.frame(colSums(is.na(df)))
```

```
##               colSums.is.na.df..
## invoice.id                0
## branch                    0
## customer.type              0
## gender                     0
## product.line               0
## unit.price                 0
## quantity                   0
## tax                        0
## date                       0
## time                       0
## payment                    0
## cogs                       0
## gross.margin.percentage    0
## gross.income               0
## rating                     0
## total                      0
```

There were no missing values.

```
#checking for duplicates  
nrow(df[duplicated(df),])
```

```
## [1] 0
```

There were no duplicates.

```
#date should be converted to datetime format  
str(df$date)
```

```
## chr [1:1000] "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" "2/8/2019" ...
```

```
#Loading the lubridate library to work with dates  
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
## hour, isoweek, mday, minute, month, quarter, second, wday, week,  
## yday, year
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## date, intersect, setdiff, union
```

```
# converting date to posixct
```

```
df$date <- as.POSIXct(df$date, format="%m/%d/%Y")  
str(df$date)
```

```
## POSIXct[1:1000], format: "2019-01-05" "2019-03-08" "2019-03-03" "2019-01-  
27" "2019-02-08" ...
```

```
# break date down to month and year and day of week components
```

```
df$month = as.factor(month(df$date, label=TRUE))
```

```
df$day = as.factor(wday(df$date, label=TRUE, week_start=1))
```

```
df$year = year(df$date)
```

```
head(df)
```

```
## invoice.id branch customer.type gender product.line unit.pric  
e  
## 1 750-67-8428 A Member Female Health and beauty 74.6  
9  
## 2 226-31-3081 C Normal Female Electronic accessories 15.2  
8  
## 3 631-41-3108 A Normal Male Home and lifestyle 46.3  
3  
## 4 123-19-1176 A Member Male Health and beauty 58.2  
2  
## 5 373-73-7910 A Normal Male Sports and travel 86.3
```

```

1
## 6 699-14-3026      C      Normal   Male Electronic accessories      85.3
9
##   quantity      tax      date  time      payment  cogs gross.margin.percen
tage
## 1          7 26.1415 2019-01-05 13:08      Ewallet 522.83      4.76
1905
## 2          5  3.8200 2019-03-08 10:29      Cash  76.40      4.76
1905
## 3          7 16.2155 2019-03-03 13:23 Credit card 324.31      4.76
1905
## 4          8 23.2880 2019-01-27 20:33      Ewallet 465.76      4.76
1905
## 5          7 30.2085 2019-02-08 10:37      Ewallet 604.17      4.76
1905
## 6          7 29.8865 2019-03-25 18:30      Ewallet 597.73      4.76
1905
##   gross.income rating      total month day year
## 1      26.1415      9.1 548.9715   Jan Sat 2019
## 2       3.8200      9.6  80.2200   Mar Fri 2019
## 3      16.2155      7.4 340.5255   Mar Sun 2019
## 4      23.2880      8.4 489.0480   Jan Sun 2019
## 5      30.2085      5.3 634.3785   Feb Fri 2019
## 6      29.8865      4.1 627.6165   Mar Mon 2019

```

#separating continuous and categorical

```

contin = c("unit.price", "quantity", "tax", "cogs", "gross.margin.percentage"
, "gross.income", "rating", "total")
cat = c("invoice.id", "branch", "customer.type", "gender", "product.line", "mo
nth", "day", "year", "payment")

```

#checking for outliers in continuous columns


#function to replace period in column names with blankspace

```

repl <- function(x){
  gsub(".", " ", x, fixed=TRUE)
}
#checking for outliers in continuous columns
for (x in contin){
  boxplot(df[x], main=repl(x), xlab=repl(x), col="blue")
}


```


unit price




unit price

quantity




quantity

tax




tax

cogs




cogs

gross margin percentage

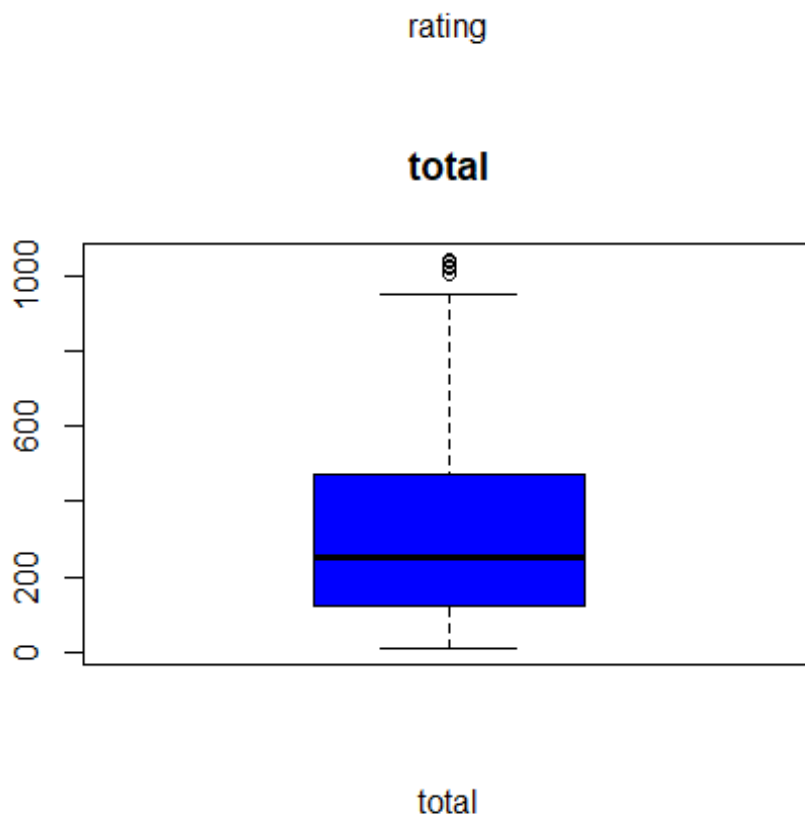



gross margin percentage

gross income



gross income



There were a few outliers in the tax, cogs(cost of goods), gross income, and total columns. They will not be dropped as it is possible for some customers to carry out shopping that costs much more than average.

```
#checking for anomalies in continuous  
#the values should not be less than zero.
```

```
for (x in contin){  
  print(paste(x, nrow(subset(df, df[x] < 0))))  
}
```

```
## [1] "unit.price 0"  
## [1] "quantity 0"  
## [1] "tax 0"  
## [1] "cogs 0"  
## [1] "gross.margin.percentage 0"  
## [1] "gross.income 0"  
## [1] "rating 0"  
## [1] "total 0"
```

```
#none of the values were less than zero
```

```
#checking for number of unique values in categorical columns
```

```
for (x in cat){  
  print(paste(x, length(unique(df[[x]]))))  
}
```

```
## [1] "invoice.id 1000"  
## [1] "branch 3"  
## [1] "customer.type 2"  
## [1] "gender 2"  
## [1] "product.line 6"  
## [1] "month 3"  
## [1] "day 7"  
## [1] "year 1"  
## [1] "payment 3"
```

```
# dropping id column because it is unique for each row, similar to the index  
df <- subset(df, select=-invoice.id)  
colnames(df)
```

```
## [1] "branch" "customer.type"  
## [3] "gender" "product.line"  
## [5] "unit.price" "quantity"  
## [7] "tax" "date"  
## [9] "time" "payment"  
## [11] "cogs" "gross.margin.percentage"  
## [13] "gross.income" "rating"  
## [15] "total" "month"  
## [17] "day" "year"
```

```
#checking for anomalies in categorical
```

```
for (x in cat[2:9]){  
  print(x)
```

```

print(unique(df[[x]]))

print("*****")
}

## [1] "branch"
## [1] "A" "C" "B"
## [1] "*****"
## [1] "customer.type"
## [1] "Member" "Normal"
## [1] "*****"
## [1] "gender"
## [1] "Female" "Male"
## [1] "*****"
## [1] "product.line"
## [1] "Health and beauty"      "Electronic accessories" "Home and lifestyle"
## [4] "Sports and travel"      "Food and beverages"     "Fashion accessories"
"
## [1] "*****"
## [1] "month"
## [1] Jan Mar Feb
## 12 Levels: Jan < Feb < Mar < Apr < May < Jun < Jul < Aug < Sep < ... < Dec
## [1] "*****"
## [1] "day"
## [1] Sat Fri Sun Mon Thu Wed Tue
## Levels: Mon < Tue < Wed < Thu < Fri < Sat < Sun
## [1] "*****"
## [1] "year"
## [1] 2019
## [1] "*****"
## [1] "payment"
## [1] "Ewallet"      "Cash"          "Credit card"
## [1] "*****"

```

No anomalous values observed. The year is 2019 only so will drop year column

```
df <- subset(df, select=-year)
```

##Univariate Analysis

```

#Loading ggplot 2 library for visualisation
library(ggplot2)

```

contin

```

## [1] "unit.price"      "quantity"
## [3] "tax"             "cogs"
## [5] "gross.margin.percentage" "gross.income"
## [7] "rating"          "total"

```

```
#statistical summary of unit price variable
```

```
data.frame(describe(df$unit.price))
```

```
##      vars      n    mean      sd median trimmed      mad   min   max range  
## X1      1 1000 55.67213 26.49463  55.23 55.6178 33.36591 10.08 99.96 89.88  
##              skew  kurtosis      se  
## X1 0.00705623 -1.222062 0.8378337
```

```
#plotting unit.price histogram
```

```
hist(df$unit.price, col="darkmagenta",  
      main="Histogram of unit price",  
      xlab="unit price")
```



In most invoices the unit price was between 90 and 100


```
#statistical summary of quantity
```

```
describe(df$quantity)
```

```
##      vars      n mean      sd median trimmed      mad min max range skew kurtosis  s  
e  
## X1      1 1000 5.51 2.92      5      5.51 2.97      1  10      9 0.01      -1.22 0.0  
9
```

```
#histogram of quantity
```

```
hist(df$quantity, col="darkmagenta",  
      main="Histogram of quantity",  
      xlab="quantity")
```



In most invoices

the quantity of units was ranging from 1 to 2

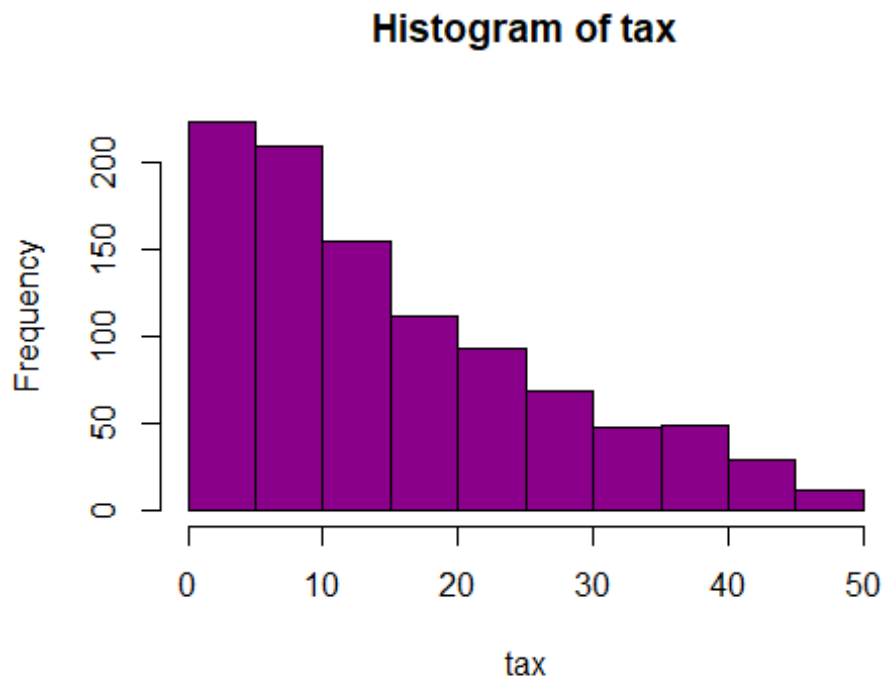
```
#statistical sumary of tax
```

```
describe(df$tax)
```

```
##      vars      n  mean    sd median trimmed   mad  min   max range skew kurtos
is
## X1      1 1000 15.38 11.71  12.09      14 11.13  0.51 49.65 49.14 0.89   -0.
09
##          se
## X1 0.37
```

```
#histogram of tax
```

```
hist(df$tax, col="darkmagenta",
      main="Histogram of tax",
      xlab="tax")
```

For most invoices, the value of tax ranged from 0 to 5.

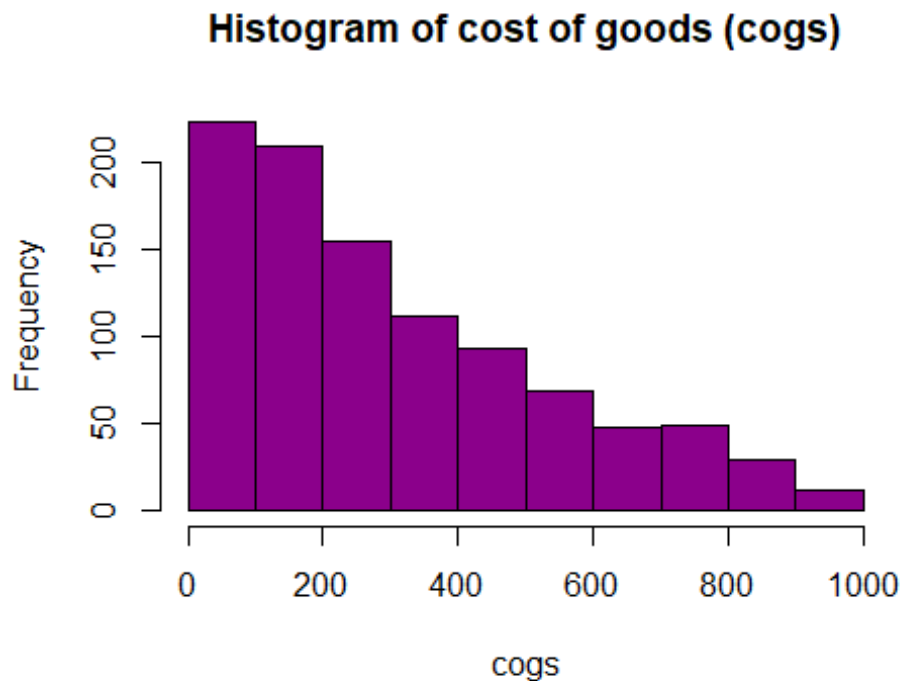
```
#statistical summary of cogs variable
```

```
describe(df$cogs)
```

```
##   vars    n  mean    sd median trimmed   mad  min max  range skew kurtosis
## X1      1 1000 307.59 234.18 241.76  279.91 222.65 10.17 993 982.83 0.89
##              -0.09
##      se
## X1 7.41
```

```
#histogram of informational_duration
```

```
hist(df$cogs, col="darkmagenta",
      main="Histogram of cost of goods (cogs)",
      xlab="cogs")
```

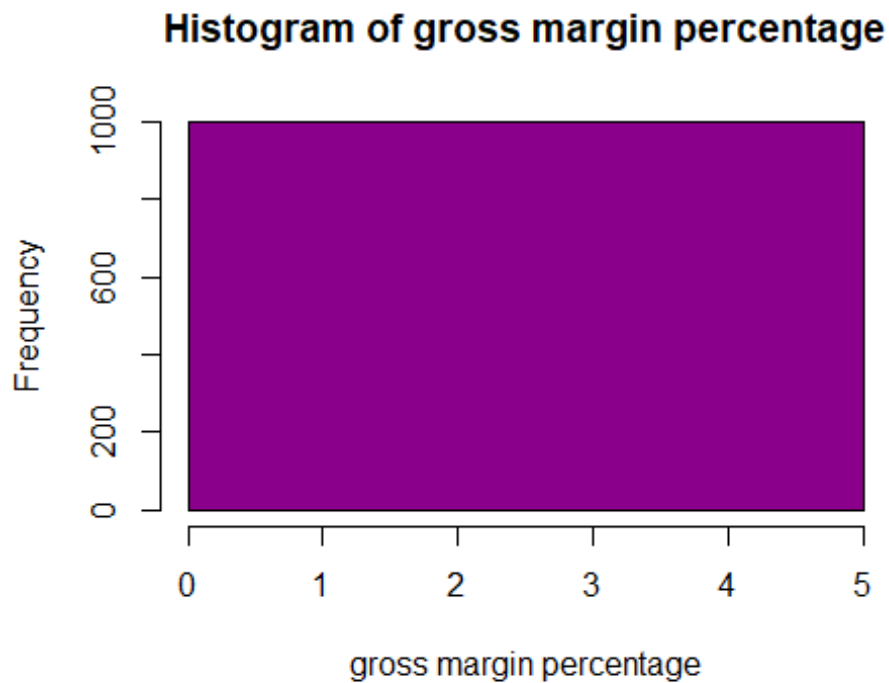


The cost of goods (without tax) in most invoices ranged from 0 to 200

```
#statistical sumary of gross margin percentage variable  
describe(df$gross.margin.percentage)
```

```
##      vars      n mean sd median trimmed mad  min  max range skew kurtosis se  
## X1      1 1000 4.76  0   4.76   4.76   0 4.76 4.76    0  NaN      NaN  0
```

```
#histogram of gross margin percentage  
hist(df$gross.margin.percentage, col="darkmagenta",  
      main="Histogram of gross margin percentage",  
      xlab="gross margin percentage")
```



All values of gross margin percentage fell between 0 and 5


```
#statistical sumary of gross income variable
```

```
describe(df$gross.income)
```

```
##      vars      n  mean    sd median trimmed   mad  min   max range skew kurtos  
is  
## X1      1 1000 15.38 11.71  12.09      14 11.13 0.51 49.65 49.14 0.89   -0.  
09  
##          se  
## X1 0.37
```

```
#histogram of gross income
```

```
hist(df$gross.income, col="darkmagenta",  
      main="Histogram of gross income",  
      xlab="gross income")
```




Gross income from most sales ranged between 0 and 10

```
#statistical sumary of rating variable
describe(df$rating)

##      vars      n mean   sd median trimmed  mad min max range skew kurtosis   s
e
## X1      1 1000 6.97 1.72      7    6.97 2.22   4  10     6 0.01    -1.16 0.0
5

#histogram of rating
hist(df$rating, col="darkmagenta",
     main="Histogram of rating",
     xlab="rating")
```




Most ratings were between 4 and 4.5

```
#statistical sumary of total variable
describe(df$total)
```

```
##      vars      n  mean      sd median trimmed   mad  min    max   range ske
w
## X1      1 1000 322.97 245.89 253.85  293.91 233.78 10.68 1042.65 1031.97 0.8
9
##      kurtosis   se
## X1      -0.09 7.78
```

```
#histogram of total
hist(df$total, col="darkmagenta",
      main="Histogram of total",
      xlab="total")
```




Most total sale values ranged between 0 and 200

```
cat
```

```
## [1] "invoice.id"      "branch"          "customer.type"  "gender"
## [5] "product.line"   "month"          "day"            "year"
## [9] "payment"

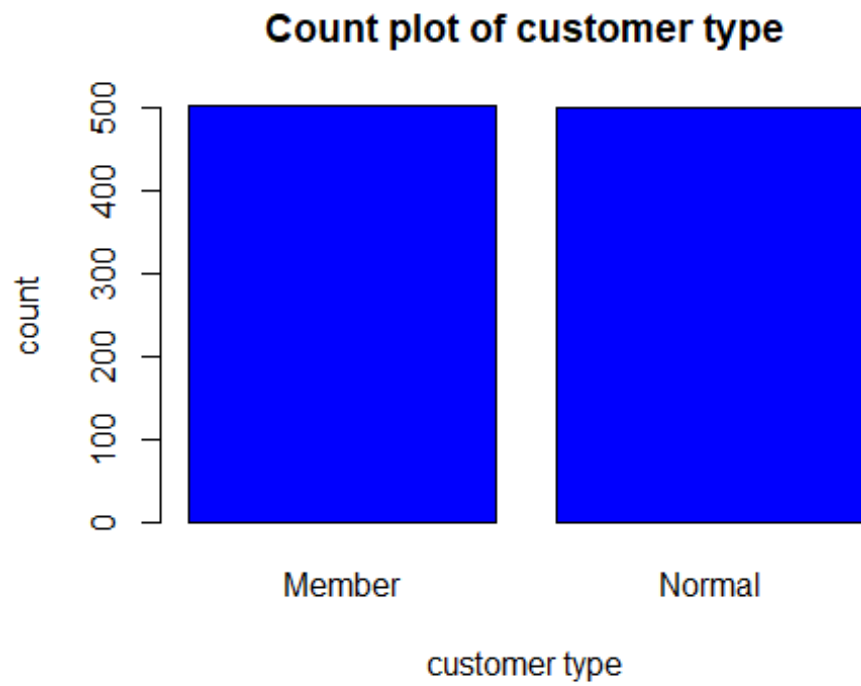
#Count plot of branch
barplot(table(df$branch), col="blue", main="Count plot of branch",
        xlab = "branch", ylab="count")
```



There is minimal difference in the number of invoices from the different branches


#count plot of customer type

```
barplot(table(df$customer.type), col="blue", main="Count plot of customer type",  
        xlab = "customer type", ylab="count")
```



The number of records from member and normal customers are almost equal

```
#count plot of gender  
barplot(table(df$gender), col="blue", main="Count plot of gender",  
        xlab = "gender", ylab="count")
```

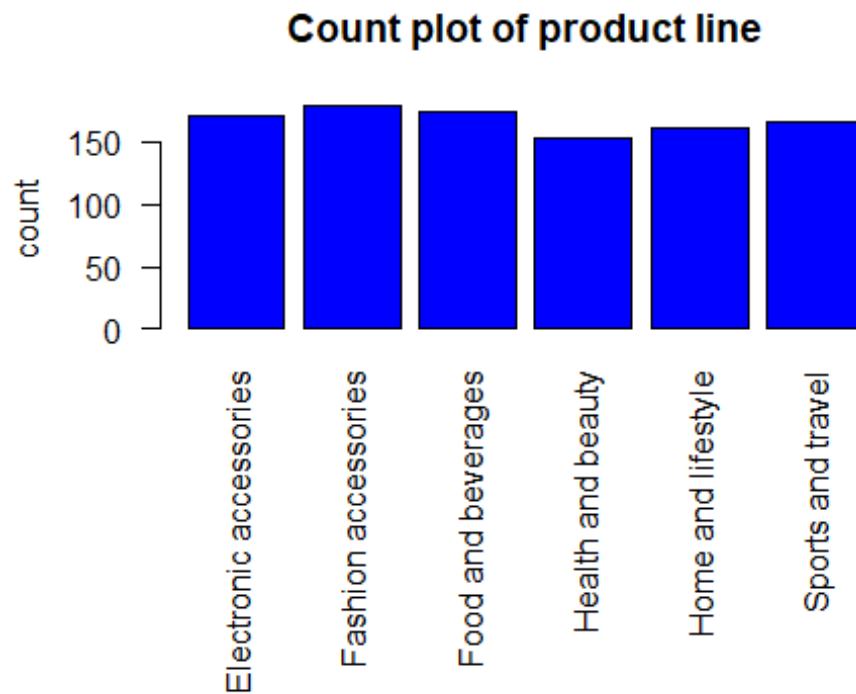



```
table(df$gender)
```

```
##  
## Female    Male  
##      501     499
```


There is minimal difference in the counts of male and female customers in the dataset.

```
#count plot of product line  
par(mar= c(10.1,4.1,4.1,2.1))  
par(las=2)  
barplot(table(df$product.line), col="blue", main="Count plot of product line"  
 , ylab="count")
```




Fashion accessories were the product line with most sale records

```
#count plot of month  
barplot(table(df$month), col="blue",  
         main="Count plot of month",  
         xlab = "month", ylab="count")
```



Between January and March 2019, January was the month with the most invoices (sales)

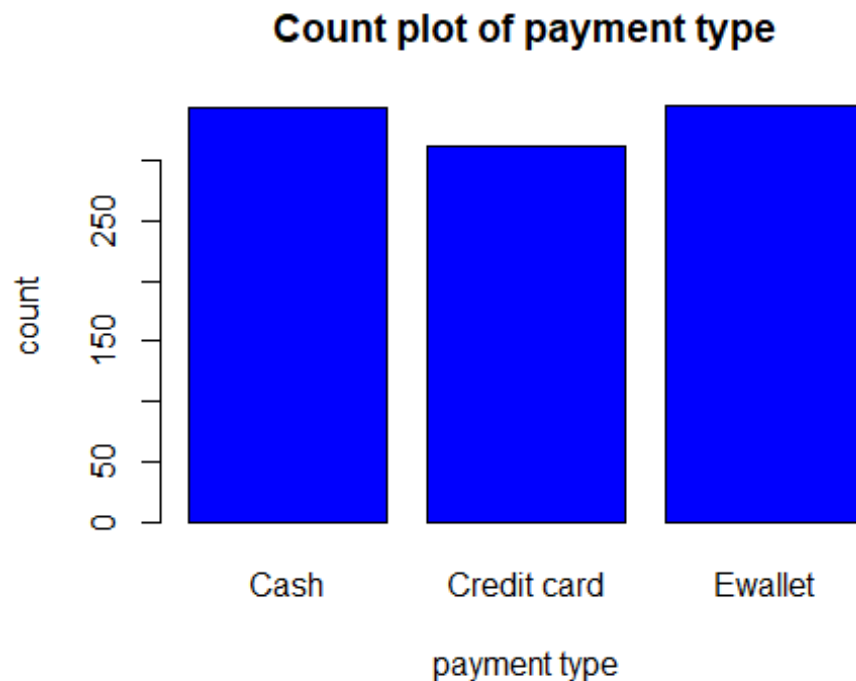
```
#count plot of day  
barplot(table(df$day), col="blue",  
        main="Count plot of day",  
        xlab = "day", ylab="count")
```



```
table(df$day)
##
## Mon Tue Wed Thu Fri Sat Sun
## 125 158 143 138 139 164 133
```

Saturday was the day with the most invoices (sales), closely followed by Tuesday.

```
#count plot of payment
barplot(table(df$payment), col="blue",
        main="Count plot of payment type",
        xlab = "payment type", ylab="count")
```



Most payments were through cash and ewallet

Bivariate Analysis

#Loading library to use functions

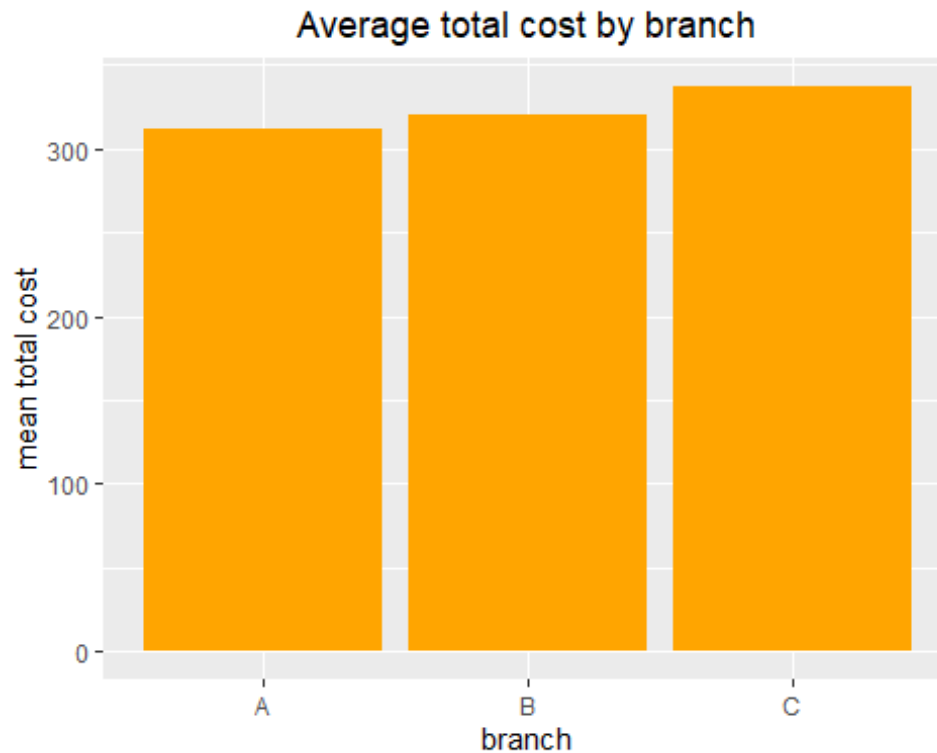
```
library("dplyr")
```

group_by

```
## function (.data, ..., .add = FALSE, .drop = group_by_drop_default(.data))
## {
##   UseMethod("group_by")
## }
## <bytecode: 0x000001fb954428c8>
## <environment: namespace:dplyr>
```

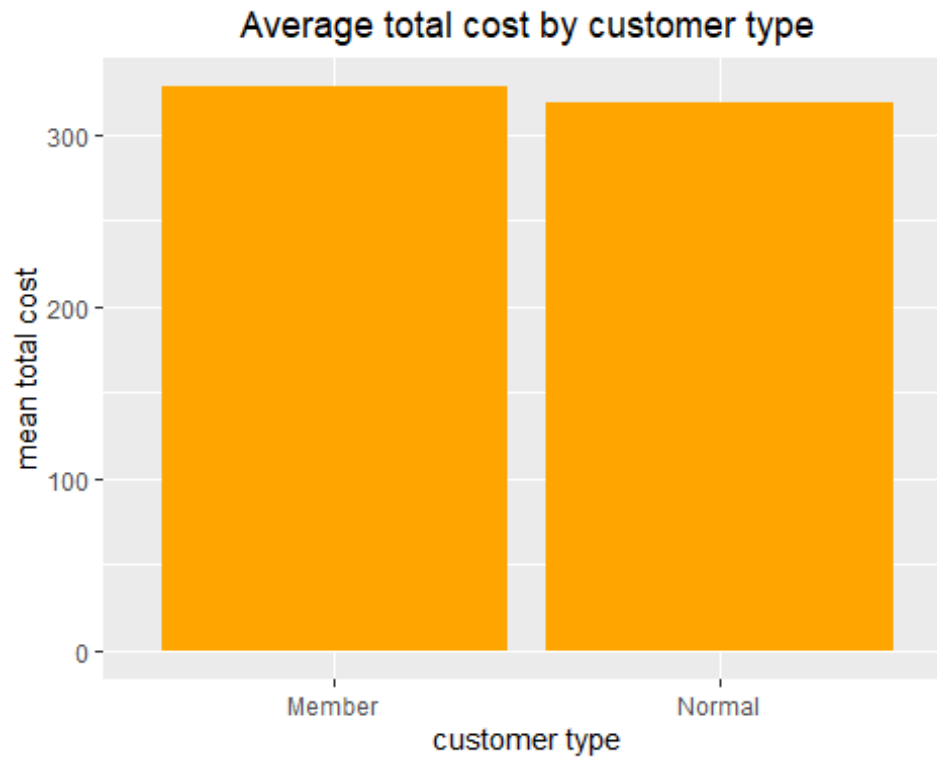
#plotting average total by branch

```
m = df %>% dplyr::group_by(branch) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(branch), y=mean),
  fill="orange") + labs(title = "Average total cost by branch",
  y="mean total cost", x="branch") + theme(plot.title =
  element_text(hjust=0.5))
```



The mean sales in branch C were higher than in the other branches


```
#plotting average total by customer type
m = df %>% dplyr::group_by(customer.type) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(customer.type), y=mean),
  fill="orange") + labs(title = "Average total cost by customer type",
  y="mean total cost", x="customer type") + theme(plot.title =
  element_text(hjust=0.5))
```



The average sales

from members were higher than from normal customers

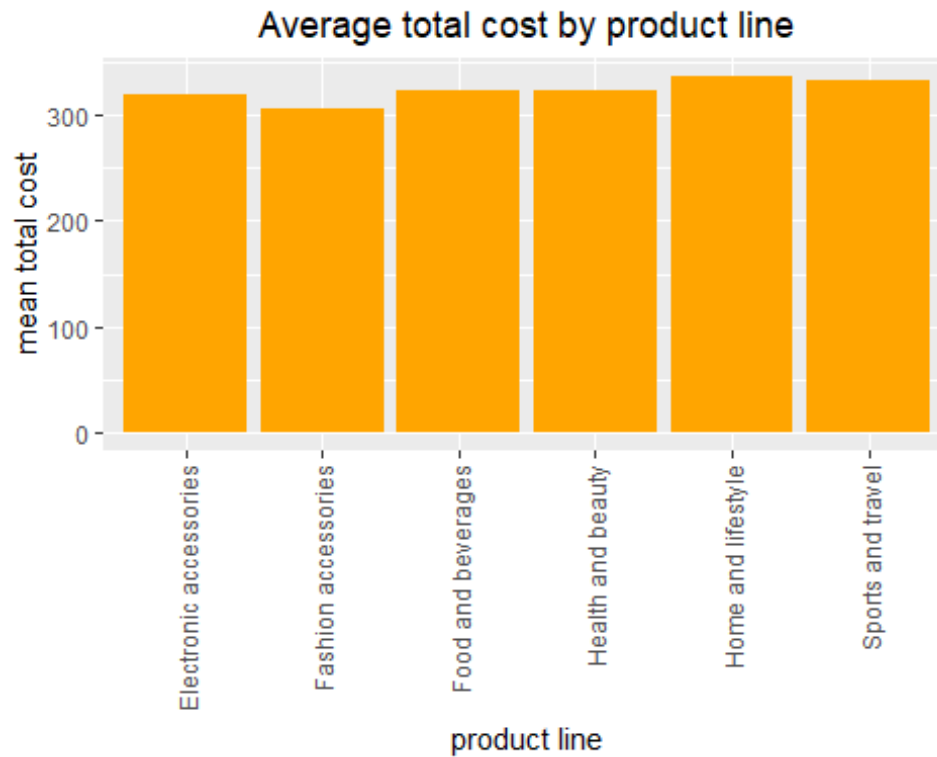
```
#plotting average total by gender
m = df %>% dplyr::group_by(gender) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(gender), y=mean),
  fill="orange") + labs(title = "Average total cost by gender",
  y="mean total cost", x="gender") + theme(plot.title =
  element_text(hjust=0.5))
```



The mean sales

were higher among female than male customers

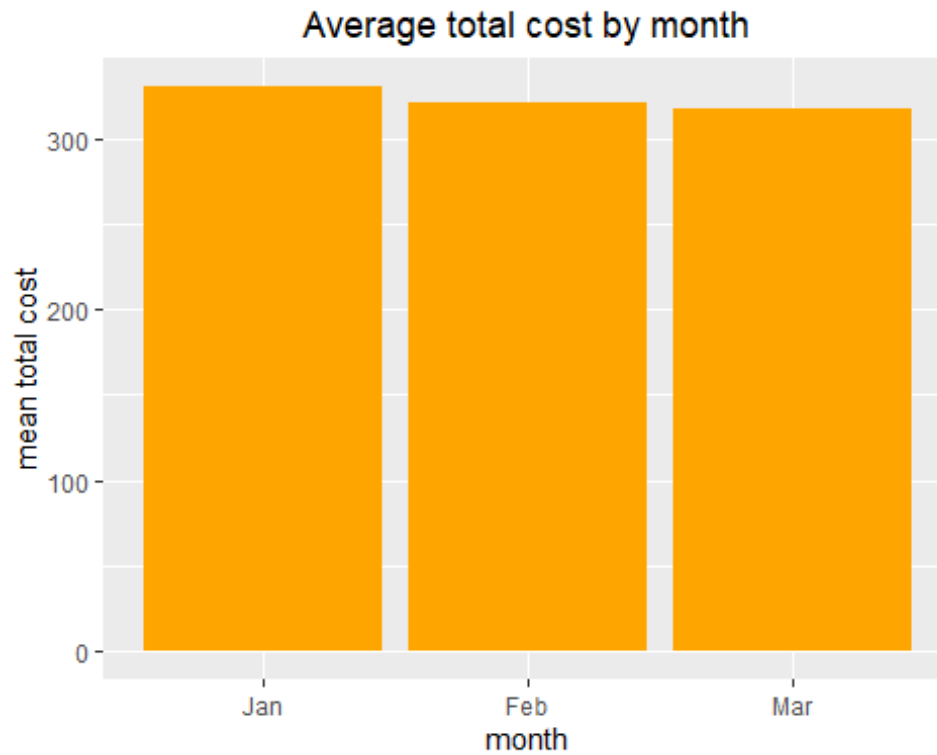
```
#plotting average total by product line
m = df %>% dplyr::group_by(product.line) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(product.line), y=mean),
  fill="orange") + labs(title = "Average total cost by product line",
  y="mean total cost", x="product line") + theme(plot.title =
  element_text(hjust=0.5), axis.text.x=element_text(angle=90,vjust=0.5,hjust=1)
)
```

The product line


with the highest average sales was home and lifestyle

```
#plotting average total by month
m = df %>% dplyr::group_by(month) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(month), y=mean),
  fill="orange") + labs(title = "Average total cost by month",
  y="mean total cost", x="month") + theme(plot.title =
  element_text(hjust=0.5))
```



January had the highest average sales


```
#plotting average total by day
m = df %>% dplyr::group_by(day) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(day), y=mean),
  fill="orange") + labs(title = "Average total cost by day",
  y="mean total cost", x="day") + theme(plot.title =
  element_text(hjust=0.5))
```



highest average sales

Saturday had the

```
#plotting average total by payment
m = df %>% dplyr::group_by(payment) %>%
  dplyr::summarise(mean=mean(total))
ggplot() + geom_col(
  data=m,
  aes(x=as.factor(payment), y=mean),
  fill="orange") + labs(title = "Average total cost by payment",
  y="mean total cost", x="payment") + theme(plot.title =
  element_text(hjust=0.5))
```



The mean sales

were almost equal across the different payment methods

Scatterplots of continuous columns

```
#continuous columns
contin

## [1] "unit.price"          "quantity"
## [3] "tax"                 "cogs"
## [5] "gross.margin.percentage" "gross.income"
## [7] "rating"              "total"

#creating dataframe that containing the continuous variables
scatterp = subset(df, select = c("unit.price", "quantity", "tax",
                                "cogs", "gross.margin.percentage", "gross.income", "rating", "total"))
head(scatterp)

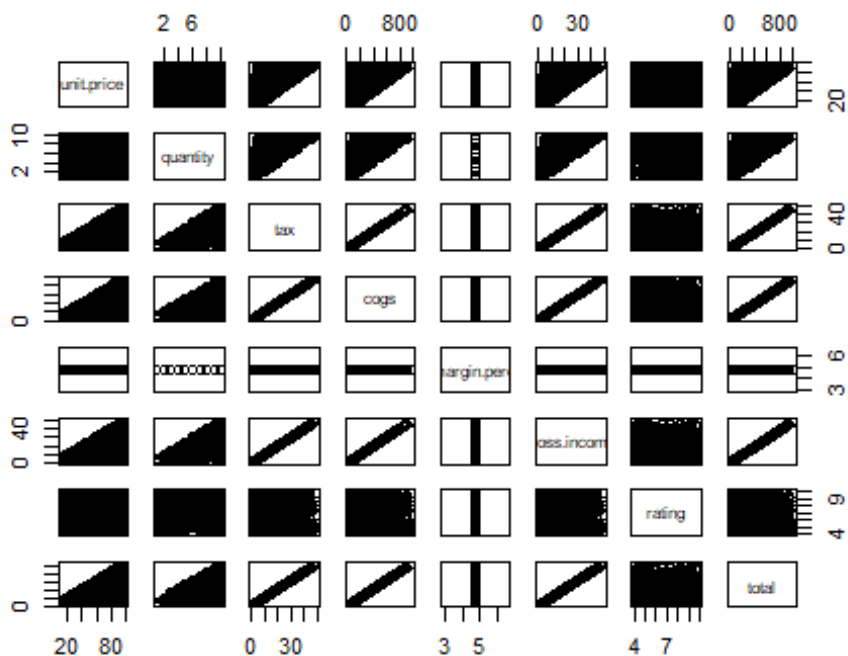
##   unit.price quantity    tax   cogs gross.margin.percentage gross.income
## 1    74.69         7 26.1415 522.83          4.761905         26.1415
## 2    15.28         5  3.8200  76.40          4.761905          3.8200
## 3    46.33         7 16.2155 324.31          4.761905         16.2155
## 4    58.22         8 23.2880 465.76          4.761905         23.2880
## 5    86.31         7 30.2085 604.17          4.761905         30.2085
## 6    85.39         7 29.8865 597.73          4.761905         29.8865
##   rating    total
## 1    9.1 548.9715
## 2    9.6  80.2200
## 3    7.4 340.5255
```

```
## 4      8.4 489.0480
## 5      5.3 634.3785
## 6      4.1 627.6165
```

```
#loading library for pair plot
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
#plotting scatterplots of continuous variables
plot(scatterp)
```



Total sales has a positive correlation with gross income, cogs(cost of goods), tax, quantity and unit price

Correlation matrix

```
str(df)
```

```
## 'data.frame':   1000 obs. of  17 variables:
## $ branch      : chr  "A" "C" "A" "A" ...
## $ customer.type : chr  "Member" "Normal" "Normal" "Member" ...
## $ gender       : chr  "Female" "Female" "Male" "Male" ...
## $ product.line : chr  "Health and beauty" "Electronic accessori
es" "Home and lifestyle" "Health and beauty" ...
## $ unit.price   : num  74.7 15.3 46.3 58.2 86.3 ...
## $ quantity     : int   7 5 7 8 7 7 6 10 2 3 ...
## $ tax          : num   26.14 3.82 16.22 23.29 30.21 ...
```

```

## $ date          : POSIXct, format: "2019-01-05" "2019-03-08" ...
## $ time          : chr  "13:08" "10:29" "13:23" "20:33" ...
## $ payment       : chr  "Ewallet" "Cash" "Credit card" "Ewallet"
...
## $ cogs          : num  522.8 76.4 324.3 465.8 604.2 ...
## $ gross.margin.percentage: num  4.76 4.76 4.76 4.76 4.76 ...
## $ gross.income   : num  26.14 3.82 16.22 23.29 30.21 ...
## $ rating         : num  9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ total          : num  549 80.2 340.5 489 634.4 ...
## $ month          : Ord.factor w/ 12 levels "Jan"<"Feb"<"Mar"<...:
1 3 3 1 2 3 2 2 1 2 ...
## $ day           : Ord.factor w/ 7 levels "Mon"<"Tue"<"Wed"<...: 6
5 7 7 5 1 1 7 4 3 ...

#apply function to time column to extract the hour of day
df$hour <- hour(as.POSIXct(df$time, format="%H:%M"))
str(df$time)

## chr [1:1000] "13:08" "10:29" "13:23" "20:33" "10:37" "18:30" "14:36" ...
str(df$hour)

## int [1:1000] 13 10 13 20 10 18 14 11 17 13 ...
unique(df$gross.margin.percentage)

## [1] 4.761905

#no variation in that column

#converting categorical to numerical
#removing date column(was broken down to components), time column(leaving the
hour aspect)
#removing gross margin percent column as there is no variation - same value
#dataframe for correlation matrix
enc_df <- subset(df, select=-c(date, time, gross.margin.percentage))

enc_df$month <- as.numeric(factor(enc_df$month))
enc_df$branch <- as.numeric(factor(enc_df$branch))
enc_df$customer.type <- as.numeric(factor(enc_df$customer.type))
enc_df$product.line <- as.numeric(factor(enc_df$product.line))
enc_df$gender <- as.numeric(factor(enc_df$gender))
enc_df$day <- as.numeric(factor(enc_df$day))
enc_df$gender <- as.numeric(factor(enc_df$gender))
enc_df$payment <- as.numeric(factor(enc_df$payment))
#checking that datatype conversion worked
str(enc_df)

## 'data.frame': 1000 obs. of 15 variables:
## $ branch : num 1 3 1 1 1 3 1 3 1 2 ...
## $ customer.type: num 1 2 2 1 2 2 1 2 1 1 ...
## $ gender : num 1 1 2 2 2 2 1 1 1 1 ...

```

```

## $ product.line : num  4 1 5 4 6 1 1 5 4 3 ...
## $ unit.price   : num  74.7 15.3 46.3 58.2 86.3 ...
## $ quantity     : int   7 5 7 8 7 7 6 10 2 3 ...
## $ tax          : num   26.14 3.82 16.22 23.29 30.21 ...
## $ payment      : num    3 1 2 3 3 3 3 3 2 2 ...
## $ cogs         : num   522.8 76.4 324.3 465.8 604.2 ...
## $ gross.income : num   26.14 3.82 16.22 23.29 30.21 ...
## $ rating       : num    9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ total        : num   549 80.2 340.5 489 634.4 ...
## $ month        : num    1 3 3 1 2 3 2 2 1 2 ...
## $ day          : num    6 5 7 7 5 1 1 7 4 3 ...
## $ hour         : int   13 10 13 20 10 18 14 11 17 13 ...

library(reshape2)

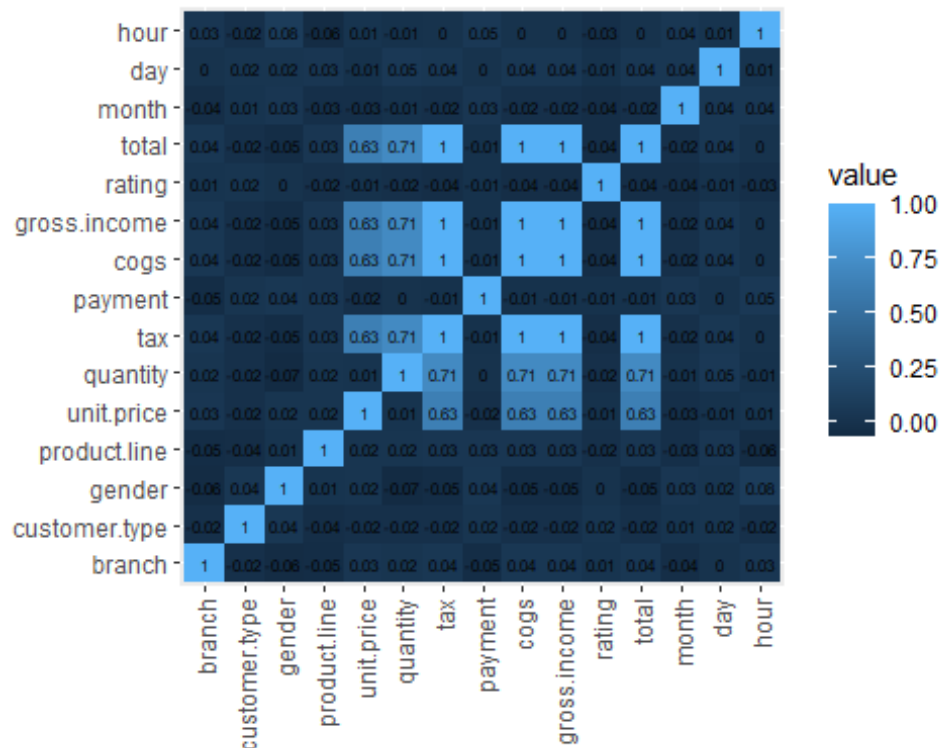
##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
##      smiths

## The following objects are masked from 'package:data.table':
##
##      dcast, melt

#plotting the correlation heatmap
datam = melt(round(cor(enc_df),2))
ggplot(data=datam, aes(x=Var1, y=Var2, fill=value)) + geom_tile() + geom_text
(aes(Var2, Var1, label=value), color="black",size=2) + theme(axis.text.x=elem
ent_text(angle=90,vjust=0.5,hjust=1), axis.title.x = element_blank(), axis.ti
tle.y = element_blank())

```



Total sales column has a perfect positive correlation with gross income, cogs (cost of goods), and tax. It also has a strong positive correlation with quantity and unit price.

##Part 1: Dimensionality Reduction

###PCA

```
# Selecting the numerical data (excluding the categorical variables)
df_num <- subset(df, select=c("unit.price", "quantity",
                              "tax", "cogs", "gross.income", "rating", "total" ))


#performing pca
df.pca <- prcomp(df_num, center = TRUE, scale. = TRUE)
summary(df.pca)

## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  2.2185 1.0002 0.9939 0.30001 2.981e-16 1.493e-16
## Proportion of Variance 0.7031 0.1429 0.1411 0.01286 0.000e+00 0.000e+00
## Cumulative Proportion 0.7031 0.8460 0.9871 1.00000 1.000e+00 1.000e+00
##              PC7
## Standard deviation  9.831e-17
## Proportion of Variance 0.000e+00
## Cumulative Proportion 1.000e+00
```

7 principal components were formed using the seven variables consisting of numerical (continuous) data. Principal component 1 explains 70.31% of the variance


```
library(factoextra)
```

```
#visualising explained variances by the principal components  
fviz_eig(df.pca)
```



Principal component 1 explains the most variance

```
# Looking at the PCA object
```

```
# ---
```

```
#
```

```
str(df.pca)
```

```
## List of 5
```

```
## $ sdev      : num [1:7] 2.22 1.00 9.94e-01 3.00e-01 2.98e-16 ...
```

```
## $ rotation: num [1:7, 1:7] -0.292 -0.325 -0.45 -0.45 -0.45 ...
```

```
## ... attr(*, "dimnames")=List of 2
```

```
## .. ..$ : chr [1:7] "unit.price" "quantity" "tax" "cogs" ...
```

```
## .. ..$ : chr [1:7] "PC1" "PC2" "PC3" "PC4" ...
```

```
## $ center   : Named num [1:7] 55.67 5.51 15.38 307.59 15.38 ...
```

```
## ... attr(*, "names")= chr [1:7] "unit.price" "quantity" "tax" "cogs" ...
```

```
## $ scale    : Named num [1:7] 26.49 2.92 11.71 234.18 11.71 ...
```

```
## ... attr(*, "names")= chr [1:7] "unit.price" "quantity" "tax" "cogs" ...
```

```
## $ x        : num [1:1000, 1:7] -2.005 2.306 -0.186 -1.504 -2.8 ...
```

```
## ... attr(*, "dimnames")=List of 2
```

```
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
```

```
## .. ..$ : chr [1:7] "PC1" "PC2" "PC3" "PC4" ...
```

```
## - attr(*, "class")= chr "prcomp"
```

```

#checking how the variables contribute to each component
#looking at the absolute values, total income, tax, cogs and gross income contribute most to pc1
df.pca$rotation

```

	PC1	PC2	PC3	PC4	PC5
unit.price	-0.29176275	0.270866890	-0.693584569	0.60037161	6.582429e-16
quantity	-0.32452880	-0.212748396	0.633152868	0.66972877	7.430508e-16
tax	-0.44977957	0.004196356	0.001836202	-0.21835146	-8.277641e-01
cogs	-0.44977957	0.004196356	0.001836202	-0.21835146	9.549992e-02
gross.income	-0.44977957	0.004196356	0.001836202	-0.21835146	2.290536e-01
rating	0.01867926	0.938775165	0.343575909	-0.01754621	-1.194541e-17
total	-0.44977957	0.004196356	0.001836202	-0.21835146	5.032105e-01

	PC6	PC7
unit.price	5.894232e-17	-7.635490e-17
quantity	1.864419e-16	-1.721827e-16
tax	1.656386e-02	2.540320e-01
cogs	-5.810190e-01	-6.350565e-01
gross.income	7.836445e-01	-2.888526e-01
rating	1.850076e-17	-7.208985e-17
total	-2.191893e-01	6.698770e-01

```
library(devtools)
```

```
## Loading required package: usethis
```


```
library(ggbiplot)
```

```
#plotting variable contributions pca
```

```


fviz_pca_var(df.pca,
  col.var = "contrib", # Color by contributions to the PC(1)?
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE        # Avoid text overlapping
)

```



total income, tax, cogs and gross income contribute most to pc1


```
#plotting the observations with using first 2 principal components  
fviz_pca_ind(df.pca,  
             label="none"  
             )
```



```
#plot of observations and variables
```

```
fviz_pca_biplot(df.pca, repel = TRUE,  
                 col.var = "#2E9FDF", # Variables color  
                 col.ind = "#696969"  # Individuals color  
                 )
```


```
## Warning: ggrepel: 973 unlabeled data points (too many overlaps). Consider  
## increasing max.overlaps
```



observations with higher values of Cogs, quantity, gross income, total, unit price and quantity are towards the left side of the plot. Most observations are towards the right side, indicating that most transactions had values on the lower end of these variables.

```
cat
## [1] "invoice.id"      "branch"          "customer.type"  "gender"
## [5] "product.line"   "month"           "day"            "year"
## [9] "payment"


#colouring observations by month
ggbiplot(df.pca,ellipse=TRUE, groups=df$month, obs.scale = 1, var.scale = 1)
```



no distinct clustering of observations. observations of the different groups are spread out in a similar manner

#colouring observations by branch


```
ggbiplot(df.pca, ellipse=TRUE, groups=df$branch, obs.scale = 1, var.scale = 1)
```



observations of the different groups are spread out in a similar manner

#colouring observations by customer type

```
ggbiplot(df.pca, ellipse=TRUE, groups=df$customer.type, obs.scale = 1, var.scale = 1)
```



the groups are spread out similarly

##Part 2: Feature Selection

####Filter method

```
library(caret)
library(corrplot)
```

corrplot 0.92 loaded

Calculating the correlation matrix

```
matrix <- cor(enc_df)
```

Finding variables that are highly correlated

findCorrelation function - removes redundancy by correlation

#

```
highlyCorrelated <- findCorrelation(matrix
                                   , cutoff=0.80)
```

Highly correlated attributes

```
highlyCorrelated
```

```
## [1] 7 9 10
```



```
names(enc_df[,highlyCorrelated])

## [1] "tax"          "cogs"         "gross.income"
```


Tax, cost of goods (cogs), and gross income have been identified as redundant features that will be dropped

```
# Removing redundant features
```

```
enc_dfc<-enc_df[-highlyCorrelated]
```

```
# plotting the correlation matrices before and after removing redundant features
```

```
par(mfrow = c(1, 2))# figure arrangement c(rows, columns)
corrplot(matrix, order = "hclust", tl.cex=0.8, cl.cex=0.5)
corrplot(cor(enc_dfc), order = "hclust", tl.cex=0.8, cl.cex=0.5)
```



```
####Embedded method
```

```
library(wskm)

## Loading required package: latticeExtra

##
## Attaching package: 'latticeExtra'
```


```
## The following object is masked from 'package:ggplot2':
##
##      layer

## Loading required package: fpc

# Model - entropy weighted k means.
# dropping the redundant columns identified above before modelling
set.seed(2)
model <- ewkm(subset(enc_df, select=-c(tax,cogs,gross.income)), 2)

library("cluster")

# cluster plot against first 2 principal components
clusplot(enc_df, model$cluster, color=TRUE, shade=TRUE,
         lines=1,main='Cluster Analysis')
```



These two components explain 40.73 % of the point variab

```
# Weights are calculated for each variable and cluster.
# They are a measure of the relative importance of each variable
# with regards to the membership of the observations to that cluster.

round(model$weights*100,2)

##   branch customer.type gender product.line unit.price quantity payment rat
ing
## 1      0          48.99  51.00           0           0           0           0
0
## 2      0          47.38  52.61           0           0           0           0
```

```
0
##   total month day hour
## 1      0      0  0    0
## 2      0      0  0    0
```

The variables that were determined to be most important in determining membership of observations to the various clusters included customer type and gender

##Part 3: Association Rules

```
# Loading the arules library
#
library(arules)

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack

##
## Attaching package: 'arules'

## The following object is masked from 'package:dplyr':
##
##   recode

## The following objects are masked from 'package:base':
##
##   abbreviate, write

# Loading the dataset from the csv file
# ---
# using read.transactions fuction which will load data from comma-separated f
iles
# and convert them to the class transactions for use in developing associatio
n rules
# ---
#
data <- "http://bit.ly/SupermarketDatasetII"

Transactions<-read.transactions(data, sep = ",")

## Warning in asMethod(object): removing duplicated items in transactions

Transactions

## transactions in sparse format with
## 7501 transactions (rows) and
## 119 items (columns)
```

```
# checking that object's class is transactions
```

```
class(Transactions)
```

```
## [1] "transactions"
```

```
## attr(,"package")
```

```
## [1] "arules"
```

```
# Previewing the first 10 transactions
```

```
#
```

```
inspect(Transactions[1:10])
```

```
##      items
```

```
## [1] {almonds,  
##      antioxydant juice,  
##      avocado,  
##      cottage cheese,  
##      energy drink,  
##      frozen smoothie,  
##      green grapes,  
##      green tea,  
##      honey,  
##      low fat yogurt,  
##      mineral water,  
##      olive oil,  
##      salad,  
##      salmon,  
##      shrimp,  
##      spinach,  
##      tomato juice,  
##      vegetables mix,  
##      whole weat flour,  
##      yams}
```

```
## [2] {burgers,  
##      eggs,  
##      meatballs}
```

```
## [3] {chutney}
```

```
## [4] {avocado,  
##      turkey}
```

```
## [5] {energy bar,  
##      green tea,  
##      milk,  
##      mineral water,  
##      whole wheat rice}
```

```
## [6] {low fat yogurt}
```

```
## [7] {french fries,  
##      whole wheat pasta}
```

```
## [8] {light cream,  
##      shallot,  
##      soup}
```

```

## [9] {frozen vegetables,
##      green tea,
##      spaghetti}
## [10] {french fries}

# previewing items in dataset
items<-as.data.frame(itemLabels(Transactions))
colnames(items) <- "Item"

# checking number of items
length(items$Item)

## [1] 119

#previewing first 10
head(items, 10)

##           Item
## 1         almonds
## 2 antioxydant juice
## 3         asparagus
## 4         avocado
## 5         babies food
## 6          bacon
## 7    barbecue sauce
## 8         black tea
## 9        blueberries
## 10        body spray

#there are 119 items in the dataset

# summary of the dataset

summary(Transactions)


## transactions as itemMatrix in sparse format with
## 7501 rows (elements/itemsets/transactions) and
## 119 columns (items) and a density of 0.03288973
##
## most frequent items:
## mineral water      eggs      spaghetti  french fries      chocolate
##          1788      1348          1306          1282          1229
##      (Other)
##          22405
##
## element (itemset/transaction) length distribution:
## sizes
##   1    2    3    4    5    6    7    8    9   10   11   12   13   14   15
## 1754 1358 1044  816  667  493  391  324  259  139  102   67   40   22   17
## 4

```

```
## 18 19 20
## 1 2 1
##
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 1.000 2.000 3.000 3.914 5.000 20.000
##
## includes extended item information - examples:
## labels
## 1 almonds
## 2 antioxidant juice
## 3 asparagus
```

Most transactions involved one item

```
# plotting the frequency of the top 15 most frequent items
itemFrequencyPlot(Transactions, topN = 15, col="darkgreen")
```



Mineral water is the most frequently occurring item in the transactions

```
# Building a model based on association rules
# using the apriori function
# support - proportion of transactions in which an itemset appears
# confidence - How often one item A appears whenever another item B appears in
a transaction. usually a conditional probability.
# lift
```

```
rules <- apriori (Transactions, parameter = list(supp = 0.01, conf = 0.8))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.8      0.1      1 none FALSE              TRUE        5    0.01      1
## maxlen target  ext
##      10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE     2     TRUE
##
## Absolute minimum support count: 75
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[119 item(s), 7501 transaction(s)] done [0.01s].
## sorting and recoding items ... [75 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [0 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

rules

## set of 0 rules
```

No rules obtained when the support is 0.01 and confidence 0.8.

```
#Decreasing the support and confidence
rulesa <- apriori (Transactions, parameter = list(supp = 0.001, conf = 0.7))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.7      0.1      1 none FALSE              TRUE        5    0.001      1
## maxlen target  ext
##      10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE     2     TRUE
##
## Absolute minimum support count: 7
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[119 item(s), 7501 transaction(s)] done [0.01s].
## sorting and recoding items ... [116 item(s)] done [0.00s].
## creating transaction tree ... done [0.01s].
## checking subsets of size 1 2 3 4 5 6 done [0.02s].
```

```
## writing ... [200 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rulesa
```

```
## set of 200 rules
```

200 rules obtained with lower support of 0.001 and confidence of 0.7

```
# summary of the model
# most rules have 4 items
# lift - A rule with a lift ~ 1 would imply that the probability of occurrence of the antecedent and that of the consequent are independent of each other.
# A rule with a lift of > 1 it would imply that those two occurrences are dependent on one another and useful for predicting in future datasets.
```

```
summary(rulesa)
```

```
## set of 200 rules
```

```
##
```

```
## rule length distribution (lhs + rhs):sizes
```

```
##   3   4   5   6
```

```
## 44 122  33   1
```

```
##
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
##   3.000  4.000  4.000  3.955  4.000  6.000
```

```
##
```

```
## summary of quality measures:
```

```
##      support      confidence      coverage      lift
```

```
## Min.   :0.001067 Min.   :0.7000 Min.   :0.001067 Min.   : 2.937
```

```
## 1st Qu.:0.001067 1st Qu.:0.7273 1st Qu.:0.001466 1st Qu.: 3.088
```

```
## Median :0.001200 Median :0.7500 Median :0.001466 Median : 3.616
```

```
## Mean   :0.001330 Mean   :0.7767 Mean   :0.001728 Mean   : 4.160
```

```
## 3rd Qu.:0.001466 3rd Qu.:0.8139 3rd Qu.:0.001866 3rd Qu.: 4.418
```

```
## Max.   :0.003066 Max.   :1.0000 Max.   :0.004133 Max.   :12.722
```

```
##      count
```

```
## Min.   : 8.00
```

```
## 1st Qu.: 8.00
```

```
## Median : 9.00
```

```
## Mean   : 9.98
```

```
## 3rd Qu.:11.00
```

```
## Max.   :23.00
```

```
##
```

```
## mining info:
```

```
##      data ntransactions support confidence
```

```
## Transactions      7501    0.001      0.7
```

```
##
```

```
##      apriori(data = Transactions, parameter = list(supp = 0.001, conf = 0.7))
```

```
call
```

```
# Ordering the rules by confidence and previewing the top 10
```



```
sorted_rules<-sort(rulesa, by="confidence", decreasing=TRUE)
```

```
inspect(sorted_rules[1:10])
```

##	lhs	rhs	support	confidence
coverage lift count				
## [1]	{french fries,			
##	mushroom cream sauce,			
##	pasta}	=> {escalope}	0.001066524	1.0000000 0.0
01066524	12.606723 8			
## [2]	{ground beef,			
##	light cream,			
##	olive oil}	=> {mineral water}	0.001199840	1.0000000 0.0
01199840	4.195190 9			
## [3]	{cake,			
##	meatballs,			
##	mineral water}	=> {milk}	0.001066524	1.0000000 0.0
01066524	7.717078 8			
## [4]	{cake,			
##	olive oil,			
##	shrimp}	=> {mineral water}	0.001199840	1.0000000 0.0
01199840	4.195190 9			
## [5]	{mushroom cream sauce,			
##	pasta}	=> {escalope}	0.002532996	0.9500000 0.0
02666311	11.976387 19			
## [6]	{red wine,			
##	soup}	=> {mineral water}	0.001866418	0.9333333 0.0
01999733	3.915511 14			
## [7]	{eggs,			
##	mineral water,			
##	pasta}	=> {shrimp}	0.001333156	0.9090909 0.0
01466471	12.722185 10			
## [8]	{herb & pepper,			
##	mineral water,			
##	rice}	=> {ground beef}	0.001333156	0.9090909 0.0
01466471	9.252498 10			
## [9]	{ground beef,			
##	pancakes,			
##	whole wheat rice}	=> {mineral water}	0.001333156	0.9090909 0.0
01466471	3.813809 10			
## [10]	{frozen vegetables,			
##	milk,			
##	spaghetti,			
##	turkey}	=> {mineral water}	0.001199840	0.9000000 0.0
01333156	3.775671 9			

#the first rule can be interpreted as: if an individual buys french fries, mushroom cream sauce and pasta, they are 100% likely to buy escalope. The remaining rules can be interpreted similarly

Examining rules related some of the top 10 frequently occurring items:

Milk

```
# ---
# Items that the customers bought before purchasing milk
# ---
#
milk <- subset(rulesa, subset = rhs %pin% "milk")

# Then order by confidence
milk<-sort(milk, by="confidence", decreasing=TRUE)
inspect(milk[1:5])

##      lhs                                rhs      support      confidence
## [1] {cake, meatballs, mineral water}    => {milk} 0.001066524 1.00000000
## [2] {escalope, hot dogs, mineral water} => {milk} 0.001066524 0.88888889
## [3] {meatballs, whole wheat pasta}      => {milk} 0.001333156 0.83333333
## [4] {black tea, frozen smoothie}        => {milk} 0.001199840 0.8181818
## [5] {burgers, ground beef, olive oil}   => {milk} 0.001066524 0.80000000
##      coverage      lift      count
## [1] 0.001066524 7.717078    8
## [2] 0.001199840 6.859625    8
## [3] 0.001599787 6.430898   10
## [4] 0.001466471 6.313973    9
## [5] 0.001333156 6.173663    8

# Items that the customers might buy after purchasing milk
# ---
#
milk <- subset(rulesa, subset = lhs %pin% "milk")

# Then order by confidence
milk<-sort(milk, by="confidence", decreasing=TRUE)
inspect(milk[1:5])

##      lhs                                rhs      support      confidence
## coverage      lift count
## [1] {frozen vegetables,
##      milk,
##      spaghetti,
##      turkey}          => {mineral water}    0.001199840  0.90000000 0.0
## 01333156 3.775671      9
## [2] {cake,
##      meatballs,
##      milk}            => {mineral water}    0.001066524  0.88888889 0.0
## 01199840 3.729058      8
## [3] {burgers,
##      milk,
##      salmon}          => {spaghetti}      0.001066524  0.88888889 0.0
## 01199840 5.105326      8
```

```
## [4] {chocolate,
##      ground beef,
##      milk,
##      mineral water,
##      spaghetti}      => {frozen vegetables} 0.001066524 0.8888889 0.0
01199840 9.325253      8
## [5] {ground beef,
##      nonfat milk}      => {mineral water}      0.001599787 0.8571429 0.0
01866418 3.595877      12
```

Ground beef:

```
# items that the customers bought before purchasing ground beef
```

```
#
groundbeef <- subset(rulesa, subset = rhs %pin% "ground beef")
```

```
# Then order by confidence
groundbeef<-sort(groundbeef, by="confidence", decreasing=TRUE)
inspect(groundbeef)
```

```
##      lhs                                rhs                support
## [1] {herb & pepper, mineral water, rice} => {ground beef} 0.001333156
## [2] {grated cheese, mineral water, rice} => {ground beef} 0.001066524
## [3] {burgers, herb & pepper, spaghetti}  => {ground beef} 0.001333156
## [4] {green tea, spaghetti, tomato sauce} => {ground beef} 0.001333156
##      confidence coverage lift count
## [1] 0.9090909 0.001466471 9.252498 10
## [2] 0.8888889 0.001199840 9.046887 8
## [3] 0.7692308 0.001733102 7.829037 10
## [4] 0.7142857 0.001866418 7.269820 10
```

```
# determining items that customers might buy if they have previously bought g
round beef
```

```
# ---
```

```
#
```

```
# Subset the rules
```

```
gbeef <- subset(rulesa, subset = lhs %pin% "ground beef")
```

```
# Ordering by confidence
```

```
gbeeft<-sort(gbeef, by="confidence", decreasing=TRUE)
```

```
# inspect top 5
```

```
inspect(gbeef[1:5])
```

```
##      lhs                                rhs                support
## [1] {green beans, ground beef}           => {spaghetti}      0.001066524
## [2] {ground beef, whole weat flour}      => {mineral water} 0.001066524
## [3] {ground beef, nonfat milk}           => {mineral water} 0.001599787
## [4] {extra dark chocolate, ground beef} => {spaghetti}      0.001466471
## [5] {green tea, ground beef, tomato sauce} => {spaghetti}      0.001333156
```

```
##      confidence coverage    lift    count
## [1] 0.7272727 0.001466471 4.177085    8
## [2] 0.7272727 0.001466471 3.051047    8
## [3] 0.8571429 0.001866418 3.595877   12
## [4] 0.7333333 0.001999733 4.211894   11
## [5] 0.8333333 0.001599787 4.786243   10
```

Eggs:

```
# items that the customers bought before purchasing eggs
```

```
# ---
```

```
#
```

```
eggs <- subset(rulesa, subset = rhs %pin% "eggs")
```

```
# Then order by confidence
```

```
eggs<-sort(eggs, by="confidence", decreasing=TRUE)
```

```
inspect(eggs)
```

```
##      lhs                                rhs    support    confidence covera
ge
## [1] {black tea, spaghetti, turkey} => {eggs} 0.001066524 0.8888889 0.0011
99840
## [2] {mineral water, pasta, shrimp} => {eggs} 0.001333156 0.8333333 0.0015
99787
## [3] {black tea, turkey}              => {eggs} 0.001466471 0.7333333 0.0019
99733
##      lift    count
## [1] 4.946258    8
## [2] 4.637117   10
## [3] 4.080663   11
```

```
# items that the customers might buy if they purchase eggs
```

```
# ---
```

```
#
```

```
eggs <- subset(rulesa, subset = lhs %pin% "eggs")
```

```
# Then order by confidence
```

```
eggs<-sort(eggs, by="confidence", decreasing=TRUE)
```

```
inspect(eggs[1:5])
```

```
##      lhs                                rhs    support    confidence    cove
rage    lift count
## [1] {eggs,
##      mineral water,
##      pasta}              => {shrimp}      0.001333156 0.9090909 0.00146
6471 12.722185    10
## [2] {brownies,
##      eggs,
##      ground beef}        => {mineral water} 0.001066524 0.8888889 0.00119
9840 3.729058    8
## [3] {chocolate,
```

```
##      eggs,
##      frozen vegetables,
##      ground beef}      => {mineral water} 0.001466471 0.8461538 0.00173
3102 3.549776 11
## [4] {chocolate,
##      eggs,
##      olive oil,
##      spaghetti}      => {mineral water} 0.001199840 0.8181818 0.00146
6471 3.432428 9
## [5] {cooking oil,
##      eggs,
##      olive oil}      => {mineral water} 0.001066524 0.8000000 0.00133
3156 3.356152 8
```

##Part 4: Anomaly Detection

```
# Load tidyverse and anomalize
# ---
#
library(tidyverse)
library(anomalize)

## == Use anomalize to improve your Forecasts by 50%! =====
## Business Science offers a 1-hour course - Lab #18: Time Series Anomaly Det
## </> Learn more at: https://university.business-science.io/p/learning-labs-pro </>

data <- fread("http://bit.ly/CarreFourSalesDataset")
data <- tibble::as_tibble(data)
```

###Checking the Data

Determining the no. of records in the dataset:

```
dim(data)

## [1] 1000 2

#the dataset has 1000 rows and 2 columns
```

Previewing the top of the dataset:

```
head(data)

## # A tibble: 6 × 2
##   Date      Sales
##   <chr>    <dbl>
## 1 1/5/2019  549.
## 2 3/8/2019   80.2
## 3 3/3/2019  341.
```

```
## 4 1/27/2019 489.  
## 5 2/8/2019 634.  
## 6 3/25/2019 628.
```

Previewing the bottom of the dataset:

```
tail(data)  
  
## # A tibble: 6 × 2  
##   Date      Sales  
##   <chr>    <dbl>  
## 1 2/18/2019 64.0  
## 2 1/29/2019 42.4  
## 3 3/2/2019 1022.  
## 4 2/9/2019 33.4  
## 5 2/22/2019 69.1  
## 6 2/18/2019 649.
```

Checking datatype of each column:

```
#date column needs to be converted from character to datetime  
str(data)  
  
## tibble [1,000 × 2] (S3: tbl_df/tbl/data.frame)  
## $ Date : chr [1:1000] "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" ...  
## $ Sales: num [1:1000] 549 80.2 340.5 489 634.4 ...  
## - attr(*, ".internal.selfref")=<externalptr>
```

Tidying the Dataset

```
#checking column names  
colnames(data)  
  
## [1] "Date" "Sales"  
  
#converting column names to lowercase  
colnames(data) = tolower(colnames(data))  
colnames(data)  
  
## [1] "date" "sales"  
  
#checking for missing values  
data.frame(colSums(is.na(data)))  
  
##      colSums.is.na.data..  
## date                   0  
## sales                   0  
  
#no missing values in any of the columns  
  
#checking for duplicates  
nrow(data[duplicated(data),])
```

```
## [1] 0

#no duplicates in the dataset

#loading the lubridate library to work with dates
library(lubridate)

# converting date to posixct
str(data$date)

## chr [1:1000] "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" "2/8/2019" ...

data$date <- as.POSIXct(data$date, format="%m/%d/%Y")
str(data$date)

## POSIXct[1:1000], format: "2019-01-05" "2019-03-08" "2019-03-03" "2019-01-
27" "2019-02-08" ...

#ordering by date in ascending order
data<-data[order(data$date, decreasing=FALSE),]
head(data, 10)

## # A tibble: 10 × 2
##   date                sales
##   <dtm>              <dbl>
## 1 2019-01-01 00:00:00  457.
## 2 2019-01-01 00:00:00  400.
## 3 2019-01-01 00:00:00  471.
## 4 2019-01-01 00:00:00  388.
## 5 2019-01-01 00:00:00  133.
## 6 2019-01-01 00:00:00  132.
## 7 2019-01-01 00:00:00  621.
## 8 2019-01-01 00:00:00  114.
## 9 2019-01-01 00:00:00  779.
## 10 2019-01-01 00:00:00  184.

#decomposing the sales column into "observed", "season", "trend", and "remain
der" columns.
td_sales <- data %>%
  time_decompose(sales)

## Converting from tbl_df to tbl_time.
## Auto-index message: index = date

## frequency = 11 seconds

## trend = 11 seconds

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo


head(td_sales)
```

```
## # A time tibble: 6 × 5
## # Index: date
##   date                observed season trend remainder
##   <dtm>                <dbl>  <dbl> <dbl>      <dbl>
## 1 2019-01-01 00:00:00    457. -11.7  455.      14.1
## 2 2019-01-01 00:00:00    400.  -5.11  420.     -14.9
## 3 2019-01-01 00:00:00    471.   6.94  384.      79.3
## 4 2019-01-01 00:00:00    388.  20.7   355.      12.2
## 5 2019-01-01 00:00:00    133. -27.2   326.    -166.
## 6 2019-01-01 00:00:00    132. -12.3   303.    -159.

# anomaly detection on the decomposed data using
# the remainder column through
# produces 3 new columns; "remainder_l1" (lower limit),
# "remainder_l2" (upper limit), and "anomaly" (Yes/No Flag).
# default alpha is 0.05
anom <- td_sales %>%
  anomalize(remainder)
head(anom[6:8])

## # A tibble: 6 × 3
##   remainder_l1 remainder_l2 anomaly
##           <dbl>         <dbl> <chr>
## 1         -912.          929. No
## 2         -912.          929. No
## 3         -912.          929. No
## 4         -912.          929. No
## 5         -912.          929. No
## 6         -912.          929. No

# creating the lower and upper bounds around the "observed" values and plotting anomalies
anom %>% time_recompose() %>%
plot_anomalies(time_recomposed = TRUE, alpha_dots = 0.5) + ggtitle("Plotting anomalies")
```

No anomalies


observed with alpha at default 0.05

```
#adjusting alpha parameter
data %>%
  time_decompose(sales) %>%
  anomalize(remainder, alpha=0.1) %>%
  time_recompose() %>%
  plot_anomalies(time_recomposed = TRUE, alpha_dots = 0.5)

## Converting from tbl_df to tbl_time.
## Auto-index message: index = date

## frequency = 11 seconds

## trend = 11 seconds
```



With alpha set at

0.1, there are some anomalous sales identified

#exploring anomalous sales

```
datar <- data %>%
  time_decompose(sales) %>%
  anomalize(remainder, alpha=0.1) %>%
  time_recompose()
```

```
## Converting from tbl_df to tbl_time.
```

```
## Auto-index message: index = date
```

```
## frequency = 11 seconds
```

```
## trend = 11 seconds
```

```
datao <- subset(filter(datar, anomaly == "Yes"), select=c(date, observed, anomaly))
```

```
datao <- data.frame(datao)
```

```
datao[order(datao$observed, decreasing=TRUE),]
```

```
##      date      observed anomaly
## 19 2019-02-15 1042.6500     Yes
## 13 2019-02-08 1039.2900     Yes
## 10 2019-01-30 1034.4600     Yes
##  4 2019-01-12 1023.7500     Yes
## 22 2019-03-02 1022.4900     Yes
## 14 2019-02-08 1020.7050     Yes
```

## 2	2019-01-09	1002.1200	Yes
## 25	2019-03-08	951.8250	Yes
## 9	2019-01-25	950.2500	Yes
## 34	2019-03-27	943.2990	Yes
## 8	2019-01-23	942.9000	Yes
## 21	2019-03-02	942.4485	Yes
## 31	2019-03-19	937.8180	Yes
## 26	2019-03-09	935.2665	Yes
## 20	2019-02-19	932.3370	Yes
## 3	2019-01-12	931.0350	Yes
## 28	2019-03-14	921.1860	Yes
## 17	2019-02-12	914.5500	Yes
## 1	2019-01-01	888.6150	Yes
## 15	2019-02-09	888.4050	Yes
## 33	2019-03-20	887.9220	Yes
## 30	2019-03-19	867.6150	Yes
## 23	2019-03-03	860.6850	Yes
## 27	2019-03-10	860.4750	Yes
## 7	2019-01-17	852.7050	Yes
## 18	2019-02-14	836.3040	Yes
## 12	2019-02-07	833.5950	Yes
## 24	2019-03-04	829.0800	Yes
## 16	2019-02-10	820.3650	Yes
## 29	2019-03-17	807.6600	Yes
## 32	2019-03-19	743.8200	Yes
## 11	2019-02-04	734.0760	Yes
## 5	2019-01-15	72.0090	Yes
## 6	2019-01-16	53.9280	Yes

##Conclusion

Following data preparation (where missing values, duplicates, outliers, column creation etc were dealt with accordingly), univariate and bivariate analysis were carried out on the first dataset providing valuable insights. Some general bivariate analysis insights include: member customers, women and Saturdays have higher average sales when compared to non-members, men and other days of the week, respectively, etc.

PCA: 7 principal components were formed using the seven variables consisting of numerical (continuous) data. Principal component 1 explains 70.31% of the variance. Observations with higher values of Cogs, quantity, gross income, total, unit price and quantity were towards the left side of the plot. Most observations are towards the right side, indicating that most transactions had values on the lower end of these variables. Colouring the observations by unique values of branch, month, and customer type did not reveal a distinct difference in how the observations belonging to the different groups were spread.

Feature selection - The first method used was filtering. Tax, cost of goods (cogs), and gross income were identified as redundant features that were dropped. - The next approach was of an embedded method. Model - entropy weighted k means. The variables that were determined to be most important in determining membership of observations to the various clusters included customer type and gender. The redundant features identified in filtering were dropped before modelling.

Association analysis: Different values of support and confidence were tested (200 rules obtained with a support of 0.001 and confidence of 0.7). The top rules by confidence were identified, and rules related to some of the top 10 frequently occurring items were examined.

Anomaly detection: Anomalous sales and their dates were identified, and ranked in descending order.

##Recommendations

Top 10 frequently bought items are mineral water, eggs, spaghetti, French fries, chocolate, green tea, milk, ground beef, frozen vegetables and pancakes. Some strategies to further promote the sales of some of these items:

- Have occasional sales(discounting the price) on these items as this would encourage people to increase the quantity they are purchasing and therefore boost overall amount of sales.
- Ground beef is often bought with spaghetti. A promotion such as buy x kgs of ground beef and get spaghetti free could boost the sales of the beef which is a high value commodity.
- A similar promotion could be done with eggs and black tea, and other item sets that have fast moving commodities.

The average sales were higher on Saturdays, therefore it is a key day for promotions to be run.

Fraud detection - some transactions on various dates were flagged as being anomalous and should therefore be investigated further to determine if they were legitimate.